


25 years of open source forecasting software

Rob J Hyndman

26 June 2025



Outline

1

R

2

Python

3

Julia

Outline

1 R

2 Python

3 Julia

Early R forecasting (c.2000)

ts package (now stats package):

- `HoltWinters()`: point forecasts only, with optional multiplicative seasonality (written by David Meyer).
- `arima()`: state space formulation of ARIMA models (written by Brian Ripley).
- `structTS()`: Basic structural models as per Harvey (written by Brian Ripley).

Early R forecasting (c.2000)

ts package (now stats package):

- `HoltWinters()`: point forecasts only, with optional multiplicative seasonality (written by David Meyer).
 - `arima()`: state space formulation of ARIMA models (written by Brian Ripley).
 - `structTS()`: Basic structural models as per Harvey (written by Brian Ripley).
- Each had a `predict()` method, but output was inconsistent.
 - `HoltWinters` did not produce prediction intervals.

forecast package for R: motivation

- Consistent output for existing methods by introducing new S3 generic `forecast()` and new S3 class `forecast`.
- New methods including `ets()`, `thetaf()`, `auto.arima()`.
- Modelling functions can be swapped while leaving code unchanged.
- Easy plotting tools with new `plot.forecast()` method.
- New forecasting tools such as `accuracy()` calculations.


forecast package for R: history

| Date | Event |
|------------------|--|
| Pre 2003 | Collection of functions used for consulting projects |
| July/August 2003 | <code>ets()</code> and <code>thetaf()</code> added |
| August 2006 | v1.0 available on CRAN |
| May 2007 | <code>auto.arima()</code> added |
| July 2008 | JSS paper (Hyndman & Khandakar) |
| September 2009 | v2.0 . Unbundled from Mcomp, fma & expsmooth |
| May 2010 | <code>arfima()</code> added |
| Feb/March 2011 | <code>tslm()</code> , <code>stlf()</code> , <code>naive()</code> , <code>snaive()</code> added |
| August 2011 | v3.0 . Box Cox transformations added |
| December 2011 | <code>tbats()</code> added |

forecast package for R: history

| Date | Event |
|---------------|--|
| April 2012 | Package moved to github |
| November 2012 | v4.0. <code>nnetar()</code> added |
| June 2013 | Major speed-up of <code>ets()</code> |
| January 2014 | v5.0. <code>tsoutliers()</code> and <code>tsclean()</code> added |
| May 2015 | v6.0. Added several new plots |
| February 2016 | v7.0. Added ggplot2 graphics & bias adjustment |
| March 2017 | v8.0. Added <code>tsCV()</code> & <code>baggedETS()</code> |
| April 2018 | v8.3. Added <code>mstl()</code> , and revised <code>auto.arima()</code> |
| April 2025 | v8.24. Last update |

forecast package for R



forecast package for R

- `auto.arima + forecast`
- `ets + forecast`
- `tbats + forecast`
- `bats + forecast`
- `arfima + forecast`
- `nnetar + forecast`
- `stlm + forecast`
- `meanf`
- `rwf, naive`
- `thetaf`
- `dshw, hw, holt, ses`
- `splinef`
- `croston`

All produce an object
of class `forecast`

forecast package for R

- `auto.arima + forecast`
- `ets + forecast`
- `tbats + forecast`
- `bats + forecast`
- `arfima + forecast`
- `nnetar + forecast`
- `stlm + forecast`
- `meanf`
- `rwf, naive`
- `thetaf`
- `dshw, hw, holt, ses`
- `splinef`
- `croston`

All produce an object
of class `forecast`

v9.0 will have new
model functions:

- `mean_model()`
- `rw_model()`
- `theta_model()`
- `spline_model()`
- `croston_model`

CRAN Task View Time Series

CRAN Task View: Time Series Analysis

Maintainer: Rob J Hyndman, Rebecca Killick

Contact: Rob.Hyndman at monash.edu

Version: 2025-05-17

URL: <https://CRAN.R-project.org/view=TimeSeries>

Source: <https://github.com/cran-task-views/TimeSeries/>


Contributions: Suggestions and improvements for this task view are very welcome and can be made through issues or pull requests on GitHub or via e-mail to the maintainer address. For further details see the [Contributing guide](#).

Citation: Rob J Hyndman, Rebecca Killick (2025). CRAN Task View: Time Series Analysis. Version 2025-05-17. URL <https://CRAN.R-project.org/view=TimeSeries>.


Installation: The packages from this task view can be installed automatically using the [ctv](#) package. For example, `ctv::install.views("TimeSeries", coreOnly = TRUE)` installs all the core packages or `ctv::update.views("TimeSeries")` installs all packages that are not yet installed and up-to-date. See the [CRAN Task View Initiative](#) for more details.

Top ten downloaded forecasting packages on CRAN


| Package | Downloads ('000) |
|------------|------------------|
| forecast | 22783 |
| bsts | 2322 |
| vars | 1851 |
| prophet | 1578 |
| smooth | 1159 |
| fabletools | 1134 |
| fable | 982 |
| greybox | 910 |
| hts | 896 |
| thief | 673 |




Top ten downloaded forecasting packages on CRAN



Top ten downloaded forecasting packages on CRAN



Top ten downloaded forecasting packages on CRAN



| Function | PIntervals | Automatic | Covariates |
|-----------------------------------|------------|-----------|------------|
| <code>stats::HoltWinters()</code> | No | No | No |
| <code>forecast::ets()</code> | Yes | Yes | No |
| <code>fable::ETS()</code> | Yes | Yes | No |
| <code>smooth::es()</code> | Yes | Yes | Yes |

forecast::ets()

```
ets(AirPassengers)
```

```
ETS(M,Ad,M)
```

```
Call:
```

```
ets(y = AirPassengers)
```

```
Smoothing parameters:
```

```
alpha = 0.7096
```

```
beta  = 0.0204
```

```
gamma = 1e-04
```

```
phi   = 0.98
```

```
Initial states:
```

```
l = 120.9939
```

```
b = 1.7705
```

```
s = 0.8944 0.7993 0.9217 1.059 1.22 1.232
```

```
1.111 0.9786 0.9804 1.011 0.8869 0.9059
```

forecast::ets()

```
ets(AirPassengers) |> forecast(h = 10)
```

| | Point Forecast | Lo 80 | Hi 80 | Lo 95 | Hi 95 |
|----------|----------------|-------|-------|-------|-------|
| Jan 1961 | 441.8 | 419.6 | 464.0 | 407.9 | 475.7 |
| Feb 1961 | 434.1 | 407.2 | 461.1 | 392.9 | 475.3 |
| Mar 1961 | 496.6 | 460.6 | 532.6 | 441.6 | 551.7 |
| Apr 1961 | 483.2 | 443.6 | 522.9 | 422.6 | 543.8 |
| May 1961 | 484.0 | 440.0 | 528.0 | 416.7 | 551.2 |
| Jun 1961 | 551.0 | 496.3 | 605.7 | 467.4 | 634.7 |
| Jul 1961 | 613.2 | 547.4 | 679.0 | 512.6 | 713.8 |
| Aug 1961 | 609.4 | 539.2 | 679.5 | 502.1 | 716.6 |
| Sep 1961 | 530.5 | 465.5 | 595.6 | 431.0 | 630.0 |
| Oct 1961 | 463.0 | 402.8 | 523.2 | 371.0 | 555.1 |

fable::ETS()

```
as_tsibble(AirPassengers) |>  
  model(ets = ETS(value)) |>  
  report()
```

Series: value

Model: ETS(M,Ad,M)

Smoothing parameters:

alpha = 0.7096

beta = 0.02041

gamma = 0.0001005

phi = 0.98

Initial states:

```
l[0] b[0] s[0] s[-1] s[-2] s[-3] s[-4] s[-5] s[-6] s[-7] s[-8] s[-  
9] s[-10]  
121 1.771 0.8944 0.7993 0.9217 1.059 1.22 1.232 1.111 0.9786 0.9804 1.011 0.8869  
s[-11]  
0.9059
```

fable::ETS()

```
as_tsibble(AirPassengers) |>  
  model(ets = ETS(value)) |>  
  forecast(h = 10)
```

```
# A fable: 10 x 4 [1M]  
# Key:      .model [1]  
  .model    index      value .mean  
  <chr>      <mth>      <dist> <dbl>  
1 ets      1961 Jan  N(442, 299)  442.  
2 ets      1961 Feb  N(434, 442)  434.  
3 ets      1961 Mar  N(497, 789)  497.  
4 ets      1961 Apr  N(483, 956)  483.  
5 ets      1961 May  N(484, 1177) 484.  
6 ets      1961 Jun  N(551, 1821) 551.  
7 ets      1961 Jul  N(613, 2636) 613.  
8 ets      1961 Aug  N(609, 2994) 609.  
9 ets      1961 Sep  N(531, 2577) 531.  
10 ets     1961 Oct  N(463, 2205) 463.
```

smooth::es()

```
es(AirPassengers)
```

Time elapsed: 1.49 seconds

Model estimated using es() function: ETS(MMdM)

With optimal initialisation

Distribution assumed in the model: Normal

Loss function type: likelihood; Loss function value: 526.8

Persistence vector g:

| alpha | beta | gamma |
|--------|--------|--------|
| 0.3536 | 0.0000 | 0.4560 |

Damping parameter: 0.9991

Sample size: 144

Number of estimated parameters: 18

Number of degrees of freedom: 126

Information criteria:

| AIC | AICc | BIC | BICc |
|------|------|------|------|
| 1090 | 1095 | 1143 | 1157 |

smooth::es()

```
es(AirPassengers) |> forecast(h = 10, interval = "parametric")
```

| | Point forecast | Lower bound (2.5%) | Upper bound (97.5%) |
|----------|----------------|--------------------|---------------------|
| Jan 1961 | 450.4 | 413.0 | 486.8 |
| Feb 1961 | 426.5 | 389.0 | 463.0 |
| Mar 1961 | 482.7 | 438.2 | 527.4 |
| Apr 1961 | 506.8 | 457.5 | 556.3 |
| May 1961 | 521.5 | 470.1 | 575.3 |
| Jun 1961 | 596.9 | 535.8 | 659.8 |
| Jul 1961 | 688.6 | 615.4 | 765.9 |
| Aug 1961 | 681.7 | 607.0 | 761.7 |
| Sep 1961 | 567.7 | 503.8 | 634.9 |
| Oct 1961 | 505.3 | 446.1 | 567.7 |

Benchmarks

```
bench::mark(  
  forecast = ets(AirPassengers) |> forecast(h = 10),  
  fable = as_tsibble(AirPassengers) |> model(ETS(value)) |> forecast(h = 10),  
  smooth = es(AirPassengers) |> forecast(h = 10, interval = "parametric"),  
  check = FALSE  
)
```

| expression | min | median | itr/sec | mem_alloc |
|------------|----------|----------|---------|-----------|
| forecast | 673.96ms | 673.96ms | 1.48 | 43.6MB |
| fable | 729.2ms | 729.2ms | 1.37 | 37.5MB |
| smooth | 1.61s | 1.61s | 0.62 | 221.9MB |

| Function | PIntervals | Automatic | Covariates |
|-------------------------------------|------------|-----------|------------|
| <code>stats::arima()</code> | Yes | No | Yes |
| <code>forecast::Arima()</code> | Yes | No | Yes |
| <code>forecast::auto.arima()</code> | Yes | Yes | Yes |
| <code>fable::ARIMA()</code> | Yes | Yes | Yes |
| <code>smooth::ssarima()</code> | Yes | No | Yes |
| <code>smooth::auto.ssarima()</code> | Yes | Yes | Yes |

forecast::auto.arima()

```
auto.arima(AirPassengers, lambda = 0)
```

Series: AirPassengers

ARIMA(0,1,1)(0,1,1)[12]

Box Cox transformation: lambda= 0

Coefficients:

| | | |
|------|--------|--------|
| | ma1 | sma1 |
| | -0.402 | -0.557 |
| s.e. | 0.090 | 0.073 |

$\sigma^2 = 0.00137$: log likelihood = 244.7

AIC=-483.4 AICc=-483.2 BIC=-474.8

forecast::auto.arima()

```
auto.arima(AirPassengers, lambda = 0) |> forecast(h = 10)
```

| | Point Forecast | Lo 80 | Hi 80 | Lo 95 | Hi 95 |
|----------|----------------|-------|-------|-------|-------|
| Jan 1961 | 450.4 | 429.5 | 472.3 | 418.9 | 484.3 |
| Feb 1961 | 425.7 | 402.8 | 449.9 | 391.2 | 463.3 |
| Mar 1961 | 479.0 | 450.1 | 509.7 | 435.6 | 526.8 |
| Apr 1961 | 492.4 | 459.9 | 527.2 | 443.5 | 546.7 |
| May 1961 | 509.1 | 472.7 | 548.2 | 454.6 | 570.0 |
| Jun 1961 | 583.3 | 538.9 | 631.5 | 516.8 | 658.5 |
| Jul 1961 | 670.0 | 615.9 | 728.9 | 589.1 | 762.1 |
| Aug 1961 | 667.1 | 610.4 | 729.1 | 582.3 | 764.2 |
| Sep 1961 | 558.2 | 508.5 | 612.8 | 484.0 | 643.8 |
| Oct 1961 | 497.2 | 451.0 | 548.1 | 428.3 | 577.2 |

fable::ARIMA()

```
as_tsibble(AirPassengers) |>  
  model(arima = ARIMA(log(value))) |>  
  report()
```

Series: value

Model: ARIMA(2,0,0)(0,1,1)[12] w/ drift

Transformation: log(value)

Coefficients:

| | ar1 | ar2 | sma1 | constant |
|------|--------|--------|---------|----------|
| | 0.5754 | 0.2614 | -0.5553 | 0.0193 |
| s.e. | 0.0843 | 0.0842 | 0.0771 | 0.0015 |

sigma^2 estimated as 0.001323: log likelihood=249.7

AIC=-489.3 AICc=-488.8 BIC=-474.9

fable::ARIMA()

```
as_tsibble(AirPassengers) |>  
  model(arima = ARIMA(log(value))) |>  
  forecast(h = 10)
```

```
# A fable: 10 x 4 [1M]  
# Key:      .model [1]  
  .model    index      value .mean  
  <chr>     <mth>      <dbl> <dbl>  
1 arima 1961 Jan t(N(6.1, 0.0013)) 453.  
2 arima 1961 Feb t(N(6.1, 0.0018)) 430.  
3 arima 1961 Mar t(N(6.2, 0.0022)) 486.  
4 arima 1961 Apr t(N(6.2, 0.0025)) 502.  
5 arima 1961 May t(N(6.3, 0.0028)) 522.  
6 arima 1961 Jun  t(N(6.4, 0.003)) 600.  
7 arima 1961 Jul  t(N(6.5, 0.0031)) 691.  
8 arima 1961 Aug  t(N(6.5, 0.0032)) 690.  
9 arima 1961 Sep  t(N(6.4, 0.0033)) 579.  
10 arima 1961 Oct t(N(6.2, 0.0034)) 516.
```

smooth::auto.ssarima()

```
auto.ssarima(log(AirPassengers))
```

Time elapsed: 2.35 seconds

Model estimated: SARIMA(0,1,3)[1](0,1,3)[12]

Matrix of MA terms:

| | Lag 1 | Lag 12 |
|-------|---------|---------|
| MA(1) | -0.4157 | -0.7397 |
| MA(2) | 0.0313 | 0.1145 |
| MA(3) | -0.1255 | 0.0747 |

Initial values were produced using backcasting.

Loss function type: likelihood; Loss function value: -287.9556

Error standard deviation: 0.0336

Sample size: 144

Number of estimated parameters: 7

Number of degrees of freedom: 137

Information criteria:

| AIC | AICc | BIC | BICc |
|--------|--------|--------|--------|
| -561.9 | -561.1 | -541.1 | -539.1 |

smooth::auto.ssarima()

```
auto.ssarima(log(AirPassengers)) |> forecast(h = 10)
```

| | Point forecast | Lower bound (2.5%) | Upper bound (97.5%) |
|----------|----------------|--------------------|---------------------|
| Jan 1961 | 6.104 | 6.039 | 6.169 |
| Feb 1961 | 6.049 | 5.973 | 6.124 |
| Mar 1961 | 6.181 | 6.096 | 6.266 |
| Apr 1961 | 6.185 | 6.094 | 6.276 |
| May 1961 | 6.224 | 6.128 | 6.321 |
| Jun 1961 | 6.372 | 6.271 | 6.474 |
| Jul 1961 | 6.506 | 6.399 | 6.612 |
| Aug 1961 | 6.512 | 6.401 | 6.623 |
| Sep 1961 | 6.324 | 6.209 | 6.440 |
| Oct 1961 | 6.201 | 6.082 | 6.321 |

Benchmarks

```
bench::mark(  
  forecast = auto.arima(AirPassengers, lambda = 0, biasadj = TRUE) |>  
    forecast(h = 12),  
  fable = as_tsibble(AirPassengers) |> model(ARIMA(log(value))) |>  
    forecast(h = 12),  
  smooth = auto.ssarima(log(AirPassengers)) |>  
    forecast(h = 12, interval="parametric"),  
  check = FALSE  
)
```

| expression | min | median | itr/sec | mem_alloc |
|------------|-------|--------|---------|-----------|
| forecast | 1.92s | 1.92s | 0.52 | 402.39MB |
| fable | 5.26s | 5.26s | 0.19 | 1.28GB |
| smooth | 3.2s | 3.2s | 0.31 | 428.36MB |

Outline

1

R

2

Python

3

Julia

Python packages with statistical models

- statsmodels (2016–2024)
- pmdarima (2017–2023)
- sktime (2019–2025)
- GluonTS (AWS, 2019–2025)
- Darts (2020–2025)
- Merlion (Salesforce, 2021–2023)
- statsforecast (Nixtla 2022–2025)
- hierarchicalforecast (Nixtla 2022–2025)
- pyhts (2022)
- aeon (2023–2025) [fork of sktime]

Most complete packages

sktime:

- AutoARIMA
- AutoETS
- BATS/TBATS
- Theta
- STLForecaster
- Croston
- Bagged-ETS
- Prophet

statsforecast:

- AutoARIMA
- AutoETS
- AutoTBATS
- Theta
- MSTL
- Croston
- TSB, ADIDA
- ARCH/GARCH

Python reconciliation packages

- hierarchicalforecast
- sktime
- pyhts
- Darts

Foundation models

- Time-LLM (Jin et al. 2023)
- TimeGPT-1 (Garza, Challu, and Mergenthaler-Canseco 2023)
- Lag-Llama (Rasul et al. 2023)
- TimesFM (Das et al. 2023)
- Tiny Time Mixers (Ekambaram et al. 2024)
- Moirai (Woo et al. 2024)
- MOMENT (Goswami et al. 2024)
- UniTS (Gao et al. 2024)
- Chronos (Ansari et al. 2024)

Other active Python forecasting software projects

- Prophet
- Neuralprophet
- Kats
- mlforecast (Nixtla 2022–2025)
- neuralforecast (Nixtla 2022–2025)

Neural forecasting methods available in NeuralForecast

“Autoformer”, “BiTCN”, “DeepAR”, “DeepNPTS”, “DilatedRNN”,
“FEDformer”, “GRU”, “HINT”, “Informer”, “iTransformer”, “KAN”,
“LSTM”, “MLP”, “MLPMultivariate”, “NBEATS”, “NBEATSx”, “NHITS”,
“NLinear”, “PatchTST”, “RNN”, “SOFTS”, “StemGNN”, “TCN”, “TFT”,
“TiDE”, “TimeMixer”, “TimeLLM”, “TimesNet”, “TSMixer”,
“TSMixerx”, “VanillaTransformer”

Outline

1

R

2

Python

3

Julia

■ forecast